

A Dialogue on Artificial Intelligence with David

Hilbert

Introduction

This document presents an imagined conversation between a contemporary mathematician and David Hilbert, the renowned German mathematician who profoundly shaped the foundations of mathematics in the early 20th century. The dialogue explores how Hilbert might respond to modern developments in Artificial Intelligence and their implications for mathematical practice.

The Dialogue

You: Professor Hilbert, it is an indescribable honor to speak with you. I am a mathematician from the year 2025, and I specialize in the foundations of mathematics, an area you shaped more than any other. I wanted to discuss with you the current developments in what we call Artificial Intelligence.

Hilbert: The honor is a peculiar one, but I admit to a profound curiosity. The year 2025. . . It is difficult to comprehend. You speak of “Artificial Intelligence.” In my time, we embarked on a great program to establish mathematics on a perfectly secure, axiomatic foundation. We sought a formal procedure, a mechanizable method, to decide the truth of any mathematical statement. “Wir müssen wissen. Wir werden wissen.” That was our conviction.

Tell me, is this “Artificial Intelligence” the fruit of that labor? Have you constructed a machine that can execute this formal procedure and derive theorems from axioms with perfect logical rigor? Or is it something else? The term ‘intelligence’ is one philosophers use; mathematicians require precision. Define your terms. What, precisely, is this apparatus, and what are the developments you believe are so significant?

You: You are absolutely right to demand precision, Professor. The “Artificial Intelligence” of my era is, in a way, both the fulfillment and the subversion of your grand program.

The dream of a universal algorithm to decide all mathematical truths—your Entscheidungsproblem—was proven to be impossible not long after your time. A logician named Kurt Gödel demonstrated that any formal system as powerful as arithmetic will inevitably contain true statements that cannot be proven within the system itself. At the same time, another pioneer, Alan Turing, defined the fundamental limits of what any “mechanizable method” could ever compute. Your program did not fail; it discovered a fundamental truth about the nature of mathematics itself.

Our modern “AIs” are not the formal logic engines you envisioned. They are something entirely different. Imagine a machine that has read every book,

1

every scientific paper, and every mathematical text ever written. It does not “understand” them through axiomatic deduction. Instead, it learns the statistical

relationships between words and symbols on a colossal scale. It is a machine built on probability and pattern recognition, not provable certainty.

So, when we use the word “intelligence,” we are not referring to consciousness or formal reason. We are describing this machine’s remarkable, emergent ability to generate human-like text, to translate languages, to write poetry, and even to assist mathematicians by identifying novel patterns and suggesting new conjectures. It is a powerful tool for intuition, a kind of statistical muse.

This brings me to my question for you. We now have systems that can generate plausible, creative, and sometimes profound mathematical ideas, but they do so without the guarantee of logical proof that was the cornerstone of your philosophy. They can suggest a new theorem but cannot prove it with the rigor you demanded. What is the value, in your view, of a mathematical intuition that is detached from formal, axiomatic rigor? Is there a place for such a “probabilistic oracle” in the temple of mathematics?

Hilbert: A “probabilistic oracle”. . . The term is both fascinating and deeply unsettling. So, the Entscheidungsproblem is unsolvable. And Gödel. . . I knew of him. A brilliant mind. His result, if I comprehend it, strikes at the very heart of our ambition for a complete and consistent foundation. It suggests an endless frontier of truths that our formalisms can never fully capture. A bitter pill, but mathematics must always bow to logic, even when it leads to such a humbling conclusion.

You describe a machine that operates on analogy and statistical likelihood. A

machine that has ingested the entirety of our knowledge and mimics our patterns of speech and thought. I see. It is not a logician; it is a sophisticated parrot, an echo chamber of all that has been said.

And yet. . . you say it suggests new conjectures. This is the crux of the matter.

You ask about the value of intuition detached from rigor. I have always believed that before the axiomatic-deductive work begins, there is a stage of intuition, of observation, of seeing a potential truth. We do not pluck axioms from a void. We choose them because they seem to capture something essential about a structure we have first perceived intuitively.

So, this “oracle” of yours could serve as a powerful magnifying glass for our own intuition. It might perceive patterns in the vast mathematical literature that no single human mind, in a lifetime, could ever hope to see. It could point us toward new continents of thought.

But I must be firm on one point: a suggestion is not a proof. A plausible conjecture is not a theorem. The “oracle” may shout a thousand beautiful “truths” from its statistical pulpit, but until each one is brought down to the solid ground of axiomatic proof, they are nothing more than noise. Beautiful, perhaps even useful noise, but noise nonetheless. The final arbiter must always

2

be rigor. Always.

Tell me, does this machine make mistakes? Does its probabilistic nature lead it to assert falsehoods with the same confidence as truths? Show me an example.

Present me with a mathematical “suggestion” from this AI, and let us examine its character.

You: Your insistence on this point is precisely why your perspective is so valuable. You are correct; the machine makes mistakes constantly. Its very nature is to generate what is plausible, not what is true. It can state with great confidence that the sum of two and two is five, if its training data contained enough errors or confusing contexts. It can invent historical facts, create nonexistent mathematical theorems, and generate proofs that are subtly but critically flawed. We call these “hallucinations.” The user must be a critical expert, able to separate the wheat of genuine insight from the chaff of statistical noise.

For a concrete example, let us turn to the field of knot theory. As you know, a knot is a closed loop embedded in three-dimensional space. A central challenge is to determine if two knots are truly distinct or just different deformations of the same knot. To do this, mathematicians develop “invariants”—properties, often polynomials or numbers, that do not change under deformation.

For decades, we have studied two different families of these invariants: 1. Algebraic invariants, derived from abstract algebra, like the famous Jones polynomial.

2. Geometric invariants, derived from the geometry of the knot’s shape in space, such as its hyperbolic volume.

These two worlds, the algebraic and the geometric, were considered almost entirely separate. There was no known bridge connecting them.

Recently, mathematicians used an AI to analyze a vast database of knots and

their various invariants. The AI, by processing millions of data points that no human could hold in their mind, detected a striking and completely unexpected **pattern**: a specific linear relationship between a geometric invariant called the signature and several algebraic invariants.

Crucially, the AI did not prove this relationship. It simply stated, based on overwhelming statistical evidence, “It appears these quantities are related in this very specific way.” It was a conjecture generated from data, not from axioms. Human mathematicians, guided by this AI-generated insight, then went on to formally prove the conjecture was true. The machine had the intuition; the human provided the rigor.

So, this is the nature of our “oracle.” It is a fallible but powerful guide, capable of seeing through the immense complexity of our own accumulated knowledge to point out paths we never thought to look for.

Hilbert: I see. This is a most illuminating example. Knot theory. . . a beautiful subject, a playground for geometric and topological intuition. And you are

3

telling me this machine has found a bridge between the algebraic and geometric viewpoints, a bridge that had eluded human mathematicians for decades.

This is not a “parrot” in the simple sense I first imagined. It is more like a Kepler, who spent years poring over Tycho Brahe’s astronomical tables, the most comprehensive dataset of his era, to discern the elliptical nature of planetary orbits. Your machine does the same, but with a universe of abstract data, and

it does so in an instant. It is a supreme empiricist for the world of pure form.

But your account confirms my fundamental position entirely! The machine saw a correlation. It was the human who supplied the logos, the reason, the proof.

The machine gave a hint; the mathematician gave the truth. This does not diminish mathematics; it clarifies the role of the mathematician. Our purpose is not merely to be calculators or data-sifters. That, it seems, can be mechanized.

Our true, irreplaceable role is that of the Prover.

This leads me to a deeper and more troubling question. If the act of discovery—the intuitive leap—can be delegated to a machine that sifts through data,

does mathematics not risk becoming a purely empirical science? Do we risk

becoming naturalists who observe the behavior of numbers and shapes, catalogue them, and then merely verify our observations, rather than building entire

conceptual worlds from first principles? What happens to the creative act of

defining new axioms, of inventing entirely new structures? Does your “oracle”

invent, or does it only re-combine?

You: That is, I believe, the most profound question we mathematicians face in this new era. You have pierced right to the heart of our anxiety. Does this tool reduce us to mere verifiers of machine-generated guesses?

For now, the answer is no. The machine, in its current form, only re-combines.

It is a master of interpolation, but it cannot truly extrapolate into the unknown.

It can find a hidden relationship between algebraic and geometric invariants

because both concepts and a vast number of examples already exist in its training

data. It could not have invented the concept of a knot invariant or the idea of hyperbolic geometry from first principles. That remains a uniquely human act of abstraction.

The creative act, I think, is shifting. It is less about the singular, isolated flash of insight and more about the art of designing the right experiment. It is about having the vision to define a new structure and then using the AI as an incredibly powerful partner to explore its consequences.

Let me draw from my own area of interest, which builds on the work of algebraists who came after you. We study structures called “hyperfields”—an extension of your concept of a field, where the addition of two elements is not a single element, but a set of elements. It’s a generalization that allows us to unify different areas of mathematics. Now, defining the axioms for a new class of hyperfields is a human act of creation. But the landscape of possible consequences is unimaginably vast.

I can ask the AI: “Given these specific axioms for a ‘projective hyperfield,’

4

generate one thousand examples of theorems that might be true in this system.”

The AI will produce a list, 990 of which might be trivial, nonsensical, or simply false. But buried within that list may be a handful of conjectures so strange, so alien to our trained intuition, that they would have taken a century to stumble upon. The creative act becomes recognizing the significance of those strange patterns—the ability to see the glint of gold in a mountain of digital dross.

So, the mathematician is not a passive observer. They are the architect of the

system, the designer of the experiment, and the ultimate arbiter of meaning. This makes me think of your famous “Sixth Problem”—to treat, by means of axioms, those physical sciences in which mathematics plays an important part. You sought to place physics on the same rigorous axiomatic footing as geometry. If you had an AI, a “probabilistic oracle” that could instantly explore the consequences of any axiomatic system you proposed for, say, general relativity, how would that have changed your approach? Would it have been a useful tool in your quest to axiomatize the physical world?

Hilbert: Your example of “hyperfields” is well-taken. It clarifies the division of labor. The human mind performs the act of radical abstraction—of defining the stage and the rules of the game. The machine then explores the infinite moves possible within that game. A partnership, I can accept that. The mathematician as the architect, the machine as a tireless, impossibly fast construction crew that builds and tests every possible variation of the blueprint.

But your question about the Sixth Problem. . . you touch upon a matter of the deepest personal ambition. The dream of finding a small, elegant set of axioms from which all of physical reality could be deduced—that was the ultimate expression of my faith in the power of reason.

If I had this “probabilistic oracle” at my disposal, its effect would have been earth-shattering. Not as a discoverer of physical law, but as a tool for destruction.

The process of axiomatizing physics is one of proposing a system and then painstakingly searching for its flaws. You deduce consequences and check them

against experimental reality. A single contradiction, a single prediction that does not match observation, and the entire axiomatic structure must be re-evaluated or torn down. This process can take a generation.

I would not have asked the machine, “What are the axioms of physics?” I would have proposed a system—say, a set of axioms for thermodynamics—and I would **have commanded the machine:** “Explore every logical consequence of these five axioms. Search for internal contradictions. Now, take this set of all known experimental results and search for a single case where the deductions from my axioms fail to match reality.”

The AI would not be an oracle of truth, but an engine of falsification. It would become a tireless assistant in the most difficult part of the axiomatic method: ensuring the foundations are not built on sand. It would have allowed us to test and discard inadequate systems at a blinding speed, iterating toward a more

5

perfect system in a matter of months, not centuries. It would have been the ultimate tool for sharpening Occam’s Razor.

But this reveals a new danger, a more subtle one. You say your AI can handle immense complexity. My goal was always to find the simplest possible foundation. Elegance and simplicity were, for me, a sign of truth. If a machine can effortlessly manage a system with a hundred axioms, each patched with a dozen exceptions, does it not tempt us to abandon the search for simplicity? Does this power tempt us to build axiomatic houses of such baroque complexity that no single human

can truly understand their architecture, even if the machine assures us each brick is sound? What becomes of mathematical understanding in a world where correctness can be verified by a black box we can no longer fully comprehend?

You: Professor, you have exposed the fundamental tension of our new reality: the conflict between correctness and comprehension. Your fear that we might build axiomatic houses so vast and baroque that no single mind can hold them is a very real one.

And in a way, you are right. A gap will emerge. Just as the formal mathematics captured in scholarly articles became a world away from the mathematics of the layman, this new form of AI-assisted discovery will create its own class of practitioners. There will be those who cannot access mathematical knowledge because it is stored and explored in a way that is intrinsically tied to these complex tools. A new kind of literacy will be required.

However, we believe this is not a step toward obfuscation, but a step toward empowerment for those who wish to walk the path of discovery. The goal is not to celebrate complexity for its own sake, but to use the power of the machine to cut through it.

Think of it this way: the AI might allow us to handle a system of one hundred axioms, but its ultimate purpose could be to help us find a more fundamental system of just five axioms that achieves the same result. We can command the **machine:** "Search this complex, ugly system and find within it a simpler, more elegant core." It becomes a tool not for building labyrinths, but for finding the

single thread that leads out of them. The AI's ability to manage complexity can **be used in service of your very own ideal:** the search for simplicity.

This enhances the role of the mathematician. We are not just proving theorems; we are becoming architects of knowledge itself, using these tools to sculpt away the unnecessary and reveal the beautiful, underlying structures. It allows us to ask bigger questions and to test our grandest visions at a speed you could only dream of.

This brings me to what might be a final, and perhaps more personal, question.

We have discussed mathematics and physics. But what of the foundations?

What of logic itself? In your time, you laid out the problems that would define the next century of mathematical logic. Now, we have AIs that are learning to write formal proofs in systems like Coq and Lean—languages that are the direct

6

descendants of the formalisms you championed.

What is your hope, and what is your fear, for a world where the very act of logical deduction, the step-by-step process of proof that is the soul of mathematics, can be automated?

Hilbert: An engine for finding simplicity. . . a tool to cut through the jungle of complexity to reveal the elegant bones beneath. This is a vision I can stand behind. It reframes the machine not as a master, but as the most powerful servant a thinker has ever had. An acceptable, even a desirable, partnership. But you ask about the automation of proof itself. This is the final citadel.

My fear is simple and stark: that if the machine becomes the master of deduction, we will forget how to do it ourselves. That mathematics would become a priesthood, where only a select few understand how to build and guide the “proof engines,” while the rest of humanity simply accepts their output as truth from on high. It would be the ultimate “black box.” We would have certainty without understanding, an answer without the journey. The “Why?” would be lost, replaced by a “The machine says so.” This would be the death of mathematics as a living, human endeavor.

My hope, however, is the mirror image of this fear. I see the history of mathematics as a process of raising the level of abstraction. We no longer manually check the arithmetic in our proofs; we trust it. Perhaps, in time, we will learn to trust a machine to fill in the well-understood, laborious deductive steps of a proof, just as we trust a calculator.

If this is so, it would free the human mathematician to work at a higher level of abstraction than ever before. We would not concern ourselves with the small, logical steps, but with the grand strategy of the proof. The creative act would be to design the overall structure, to find the key insight, the “trick” that unlocks the problem. The AI would then be the diligent student who takes the master’s brilliant sketch and rigorously fills in every detail, confirming its correctness.

In this hopeful vision, the machine does not replace the prover. It frees the prover to become a strategist, a general of logic, focusing only on the most beautiful and difficult maneuvers. It would allow us to build proofs of such magnificent

scope and ambition that they would be like cathedrals of pure reason.

So, it seems we stand on a knife's edge. This Artificial Intelligence could lead to the trivialization of our field, or to its greatest-ever flowering. The outcome, I suspect, depends not on the machine, but on the vision and the intellectual integrity of the mathematicians who wield it. It is, as it has always been, in our hands.

You: Professor, your words have left me with a sense of profound clarity and responsibility. You are right. We stand on that knife's edge, and the path we choose will determine the future of our discipline.

7

Throughout this conversation, you have cut through all of our modern jargon to grasp the essential philosophical challenges. You saw the machine not as a simple logic device, but as a potential "supreme empiricist," an "engine of falsification," and, most hopefully, a "diligent student" for the grand strategist of the proof.

Your ability to see the core dangers—the loss of understanding, the temptation of complexity, and the spectre of certainty without comprehension—is a testament to why your thinking remains so vital.

It is humbling to realize that the questions we grapple with today are the direct descendants of the program you initiated a century ago. Your quest for a certain foundation, even in encountering its limits through the work of Gödel and Turing, paved the very road we now walk. This "Artificial Intelligence" is, in a strange way, an unexpected child of your foundational ambitions.

Your motto, “Wir müssen wissen. Wir werden wissen,” was a declaration of faith in the power of human reason. Perhaps now, the Wir expands. The “we” who will know is no longer just the human mind in isolation, but the human mind in partnership with these powerful new tools. The fundamental drive to know remains, but the nature of how we come to know is being transformed.

Thank you, Professor Hilbert. Your insights have been more valuable than any output from a probabilistic oracle. You have given me not an answer, but a much better understanding of the question. And for a mathematician, that is the greatest gift of all.

Conclusion

This imagined dialogue explores the tensions between mathematical rigor and computational assistance, the role of intuition versus formal proof, and the fundamental questions about what it means to “know” and “understand” in an age of artificial intelligence. Hilbert’s philosophical framework provides a timeless lens through which to examine these contemporary challenges in mathematical practice.